

DEEP LEARNING WITH KERAS


Presented by

Bharathi Raja Asoka Chakravarthi

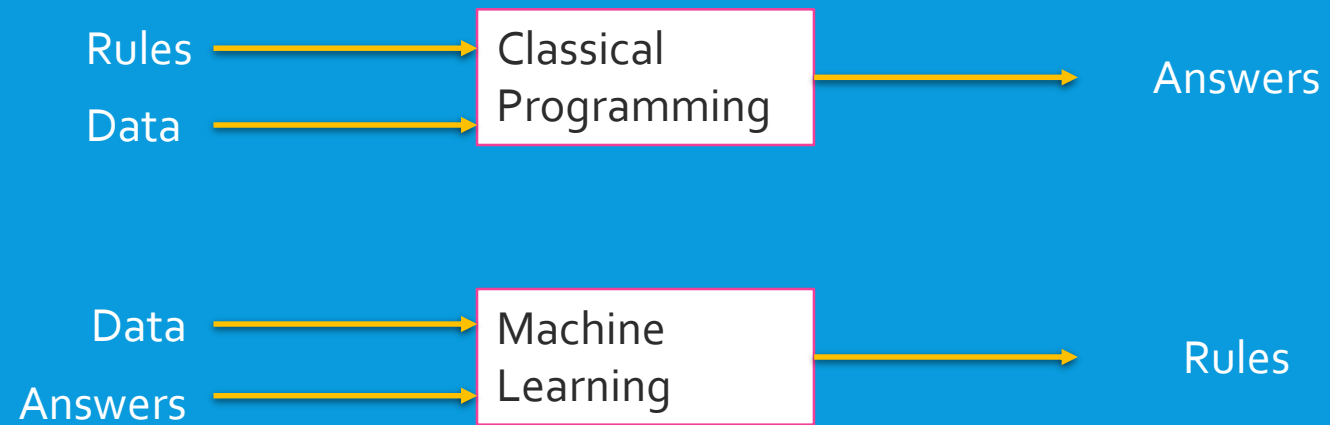
OUTLINE

- Introduction: What is deep learning?
- The mathematical building blocks of neural network
- Fundamentals of machine learning
- Deep learning for text and sequences
- Advanced deep learning practices
- Practical session: Bharathi and Bernardo


WHAT IS DEEP LEARNING?



MACHINE LEARNING

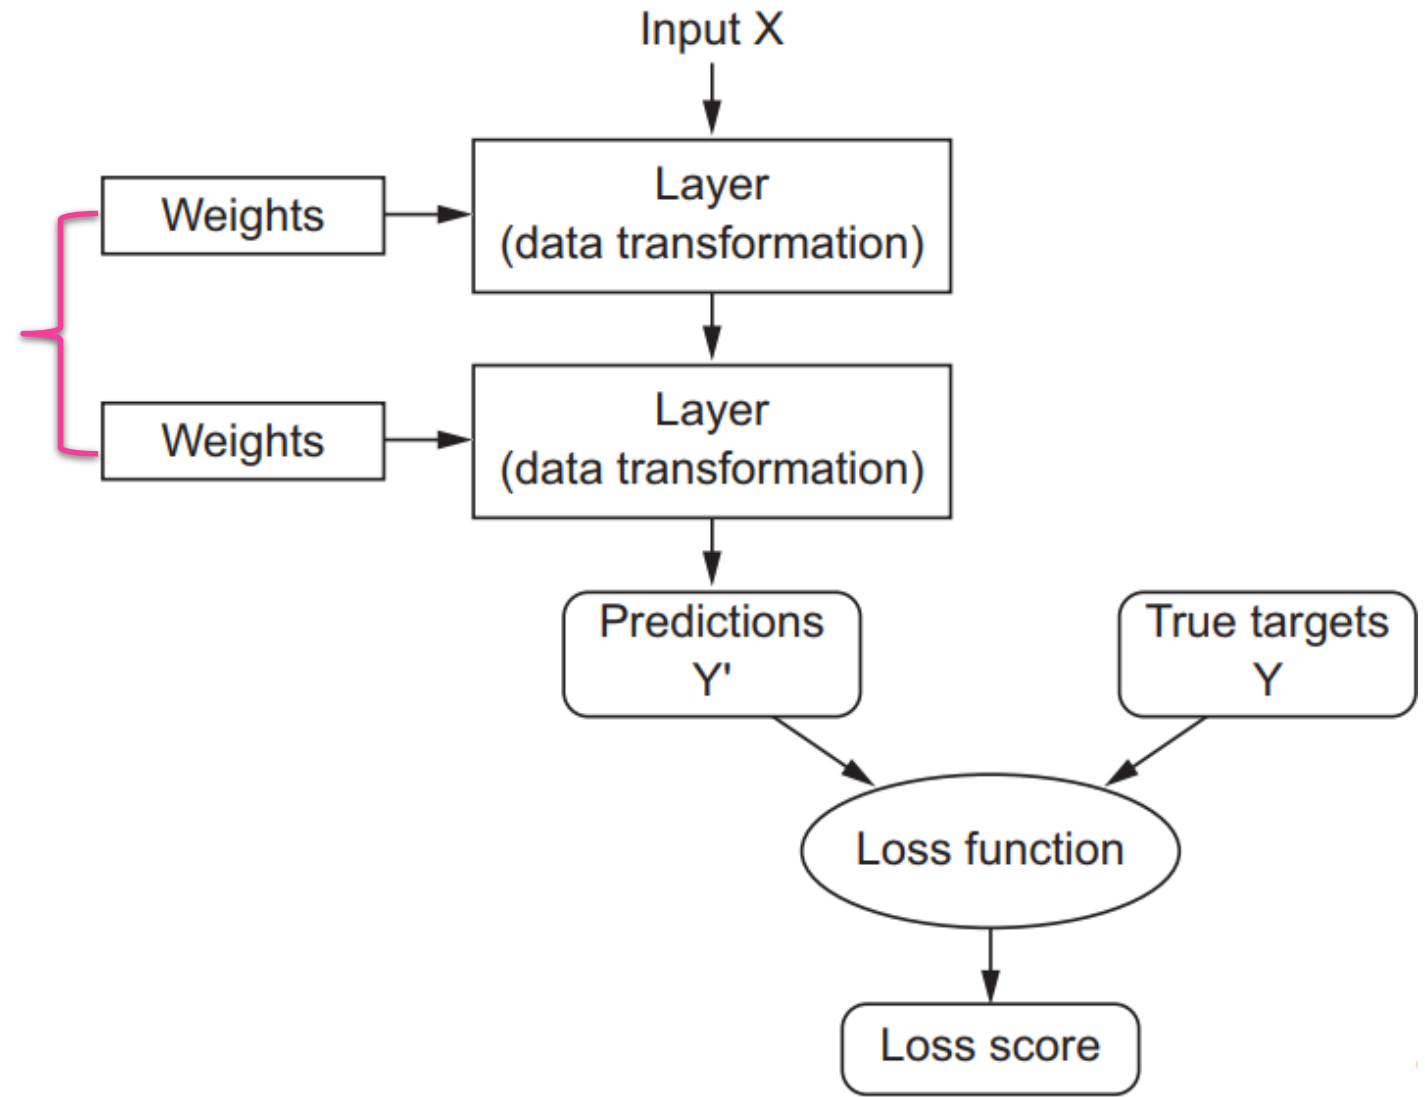


DEEP LEARNING




DEEP LEARNING

- Goal: Finding the right values for the weights



TRADITIONAL MACHINE LEARNING



DEEP LEARNING



MATHEMATICAL BUILDING BLOCKS OF NEURAL NETWORK

- A tensor is a generalization of vectors and matrices to potentially higher dimensions.


- Scalar (0D tensors)



- Vectors (1D tensors)




- Matrices (2D tensors)



Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




Multi-Class Classification with NN and SoftMax Function




GRADIENT DESCENT



EVALUATION MACHINE LEARNING MODELS




OVERFITTING AND UNDERFITTING




PREVENT OVERFITTING

- Get more training data
- Reducing the network size
- Regularization
 - L1 regularization—The cost added is proportional to the absolute value of the weight coefficients
 - L2 regularization—The cost added is proportional to the square of the value of the weight coefficients
- Dropout



(a) Standard Neural Net




(b) After applying dropout.

AVOID UNDERFITTING

- Increase model complexity
- Reduce the regularization parameters

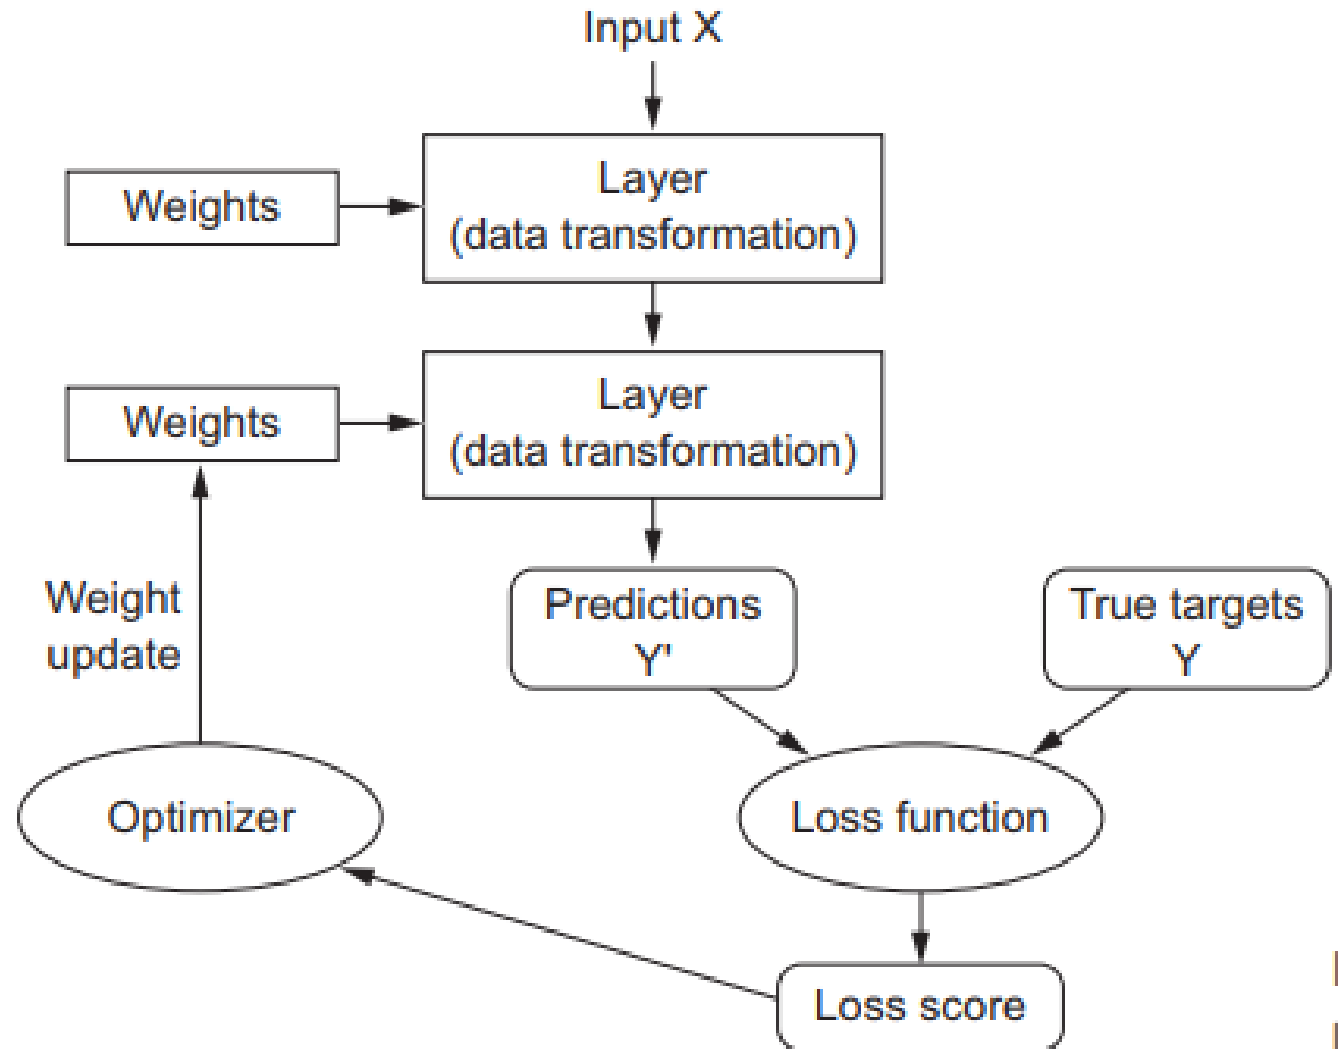
NEURAL NETWORK

- Input layer
- Hidden layer
- Output layer



ANATOMY OF A NEURAL NETWORK

- Layers, which are combined into a network (or model)
- The input data and corresponding targets
- The loss function, which defines the feedback signal used for learning
- The optimizer, which determines how learning proceeds



KERAS

- Keras is multi-backend, multi-platform
- Easy productization of models

Keras API

TensorFlow / CNTK / MXNet / Theano / ...

GPU

CPU


TPU

THREE API STYLES

- The sequential model
- The functional API
- Model sub classing


THE SEQUENTIAL API

```
import keras
from keras import layers
model=keras.Sequential()
model.add(layers.Dense(4,activation='relu',input_shape=(3,)))
model.add(layers.Dense(4,activation='relu'))
model.add(layers.Dense(10,activation='softmax'))
model.fit(x,y, epochs=10, batch_size=32)
```



THE FUNCTIONAL API

```
import keras  
from keras import layers, Input  
inputs=Input(shape=(3,))  
x=layers.Dense(4,activation='relu')(inputs)  
x=layers.Dense(4,activation='relu')(x)  
outputs=layers.Dense(10,activation='softmax')(x)  
model=keras.Model(inputs,outputs)  
model.fit(x,y,epochs=10,batch_size=32)
```



MODEL SUB CLASSING

```
import keras
```

```
from keras import layers
```

```
class MyModel(keras.Model):
```

```
    def __init__(self):
```

```
        super(MyModel, self).__init__()
```

```
        self.dense1=layers.Dense(4,activation='relu')
```

```
        self.dense2=layers.Dense(4,activation='relu')
```

```
        self.dense3=layers.Dense(10,activation='softmax')
```

```
    def call(self, inputs):
```


```
        x=self.dense1(x)
```

```
        x=self.dense2(x)
```

```
        return self.dense3(x)
```

```
model =MyModel()
```

```
model.fit(x,y, epochs=10,batch_size=32)
```



DEEP LEARNING FOR TEXT AND SEQUENCES

The cat sat on the mat



Vectors: number representation



```
from keras.preprocessing.text import Tokenizer
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
tokenizer = Tokenizer(num_words=1000)
tokenizer.fit_on_texts(samples)
sequences = tokenizer.texts_to_sequences(samples)
one_hot_results = tokenizer.texts_to_matrix(samples, mode='binary')
word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
```

USING WORD EMBEDDING

- One-hot word vectors
 - Sparse
 - High-dimensional
 - Hardcoded
 - N: total number of words (1M, 2M
- Word embeddings
 - Dense
 - Lower-dimensional
 - Learned from data
 - K: dimension (say 100,200,..)

One-hot word vectors

	1	2	3	4	5	...	N
Man	1	0	0	0	0	...	0
Girl	0	1	0	0	0	...	0
Boy	0	0	1	0	0	...	0
....							

Word embeddings

	1	2	3	4	5	...	K
Man	0.86	0.02	0.05	0.02	0.10	...	0.05
Girl	0.23	0.65	0.03	0.04	0.22	...	0.03
Boy	0.03	0.08	0.95	0.01	0.02	...	0.04
....							


LEARNING WORD EMBEDDING WITH THE EMBEDDING LAYER

Word index → Embedding layer → Corresponding word vector

```
from keras.layers import Embedding  
embedding_layer = Embedding(1000, 64)
```

RECURRENT NEURAL NETWORK

- It processes sequences by iterating through the sequence elements and maintaining a state containing information relative to what it has seen so far.



SIMPLE RNN IN KERAS

```
from keras.models import Sequential
from keras.layers import Embedding, SimpleRNN

model = Sequential()
model.add(Embedding(10000, 32))
model.add(SimpleRNN(32))
model.summary()
```

CLASSIFICATION OF IMDB MOVIE REVIEW

- `from keras.datasets import imdb`
- `from keras.preprocessing import sequence`
- `max_features = 10000`
- `maxlen = 500`
- `batch_size = 32`
- `(input_train, y_train), (input_test, y_test) =
imdb.load_data(num_words=max_features)`
- `input_train = sequence.pad_sequences(input_train, maxlen=maxlen)`
- `input_test = sequence.pad_sequences(input_test, maxlen=maxlen)`

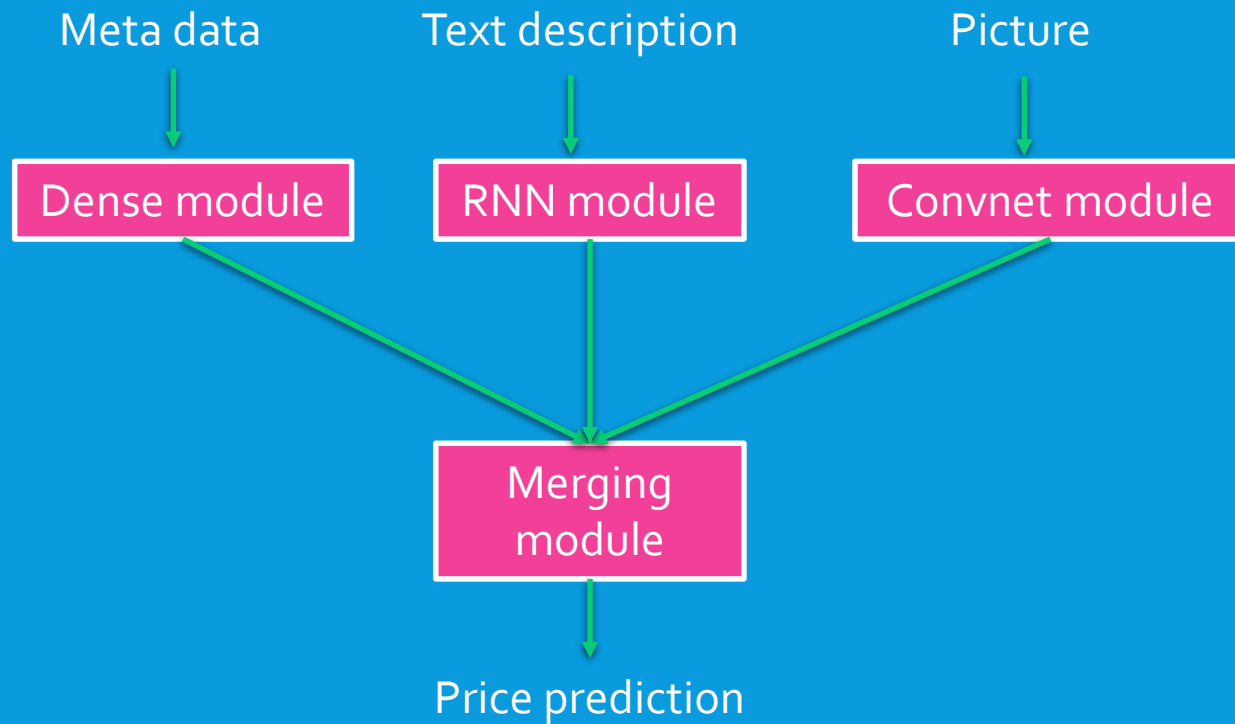
TRAINING RNN

- `from keras.layers import Dense`
- `model = Sequential()`
- `model.add(Embedding(max_features, 32))`
- `model.add(SimpleRNN(32))`
- `model.add(Dense(1, activation='sigmoid'))`
- `model.compile(optimizer='rmsprop',
loss='binary_crossentropy', metrics=['acc'])`
- `history = model.fit(input_train, y_train, epochs=10,
batch_size=128, validation_split=0.2)`

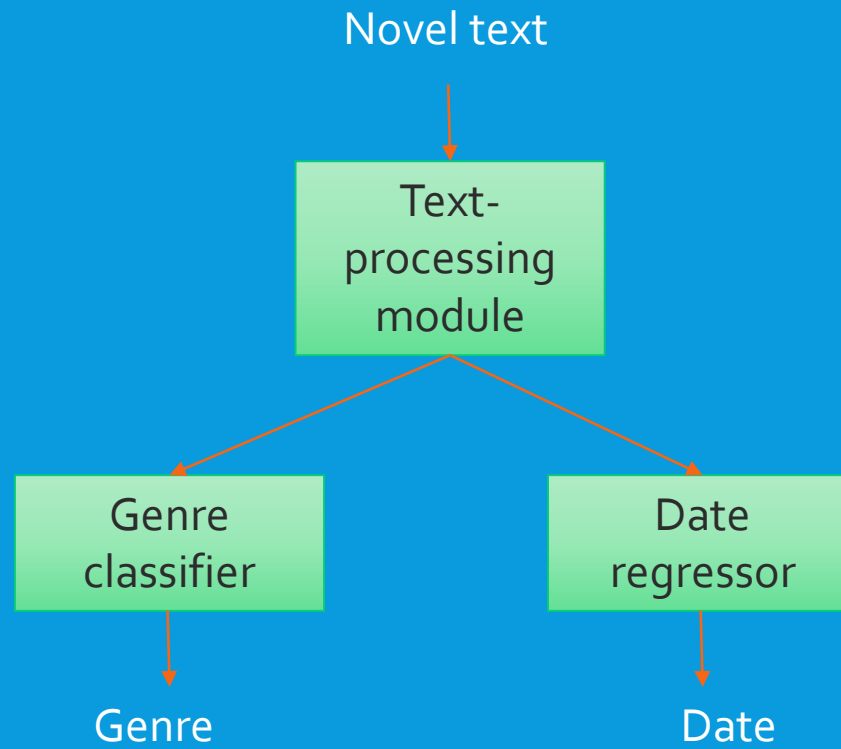
ADVANCED DEEP LEARNING PRACTICES

- Multi-input model
- Multi-output model

MULTI-INPUT MODEL



MULTIPLE OUTPUT MODEL



REFERENCES

- Francois Chollet. 2017. Deep Learning with Python (1st ed.). Manning Publications Co., Greenwich, CT, USA.
- <https://ml-cheatsheet.readthedocs.io/en/latest/optimizers.html>
- <https://keras.io/>
- <https://www.tensorflow.org/guide/keras>



THANKS